

Chapter 1

Looking at Data—

Distributions

Introduction

1.1 Data

1.2 Displaying Distributions with Graphs

1.3 Describing Distributions with Numbers

1.4 Density Curves and Normal Distributions

1.1 Data

- Cases, labels, variables, and values
- Categorical and quantitative variables

Cases, Labels, Variables, and Values

- ✓ **Cases** are the objects described by a set of data. Cases may be customers, companies, subjects in a study, units in an experiment, or other objects.
- ✓ A **label** is a special variable used in some data sets to distinguish the different cases.
- ✓ A **variable** is a special characteristic of a case.
- ✓ Different cases can have different **values** of a variable.

Categorical and Quantitative Variables

- A **categorical** variable places each case into one of several groups, or categories.
- A **quantitative** variable takes numerical values for which arithmetic operations such as adding and averaging make sense.

1.2 Displaying Distributions with Graphs

- Exploratory Data Analysis
- Graphs for categorical variables
 - Bar graphs
 - Pie charts
- Graphs for quantitative variables
 - Histograms
 - Stemplots

Exploratory Data Analysis

Exploring Data

- Begin by examining each variable by itself. Then move on to study the relationships among the variables.
- Begin with a graph or graphs. Then add numerical summaries of specific aspects of the data.

Distribution of a Variable

7

To examine a single variable, we graphically display its **distribution**.

- The distribution of a variable tells us what values it takes and how often it takes these values.
- Distributions can be displayed using a variety of graphical tools. The proper choice of graph depends on the nature of the variable.

Categorical variable

Pie chart
Bar graph

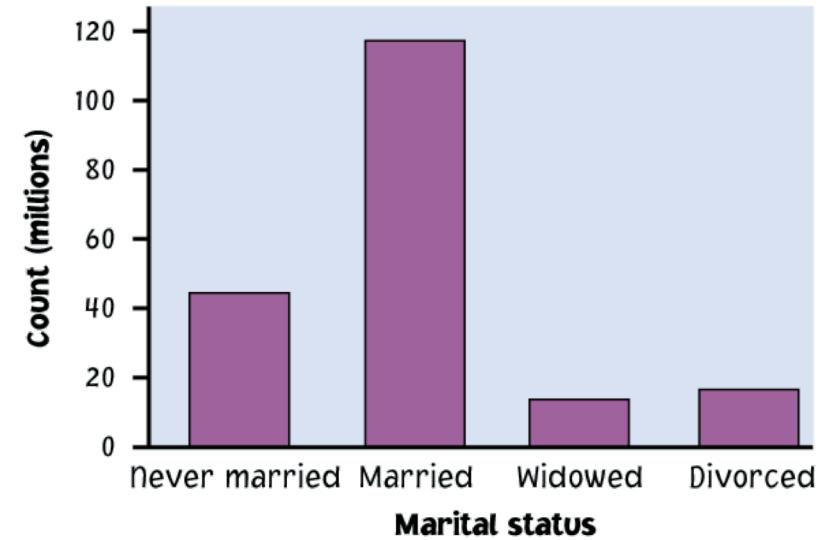
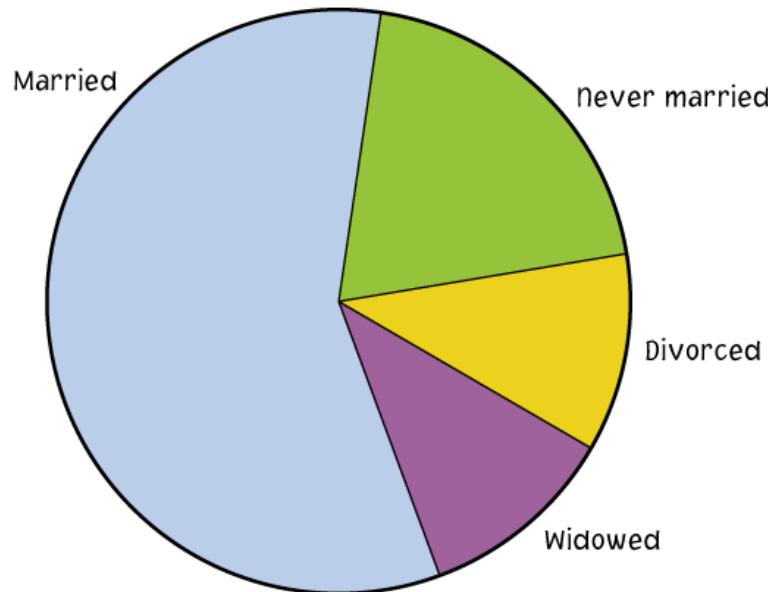
Quantitative variable

Histogram
Stemplot

Categorical Variables

The **distribution of a categorical variable** lists the categories and gives the **count** or **percent** of individuals who fall into each category.

- **Pie charts** show the distribution of a categorical variable as a “pie” whose slices are sized by the counts or percents for the categories.
- **Bar graphs** represent categories as bars whose heights show the category counts or percents.



Quantitative Variables

The **distribution of a quantitative variable** tells us what values the variable takes on and how often it takes those values.

- **Stemplots** separate each observation into a stem and a leaf that are then plotted to display the distribution while maintaining the original values of the variable.
- **Histograms** show the distribution of a quantitative variable by using bars. The height of a bar represents the number of individuals whose values fall within the corresponding class.

Stemplots 1

10

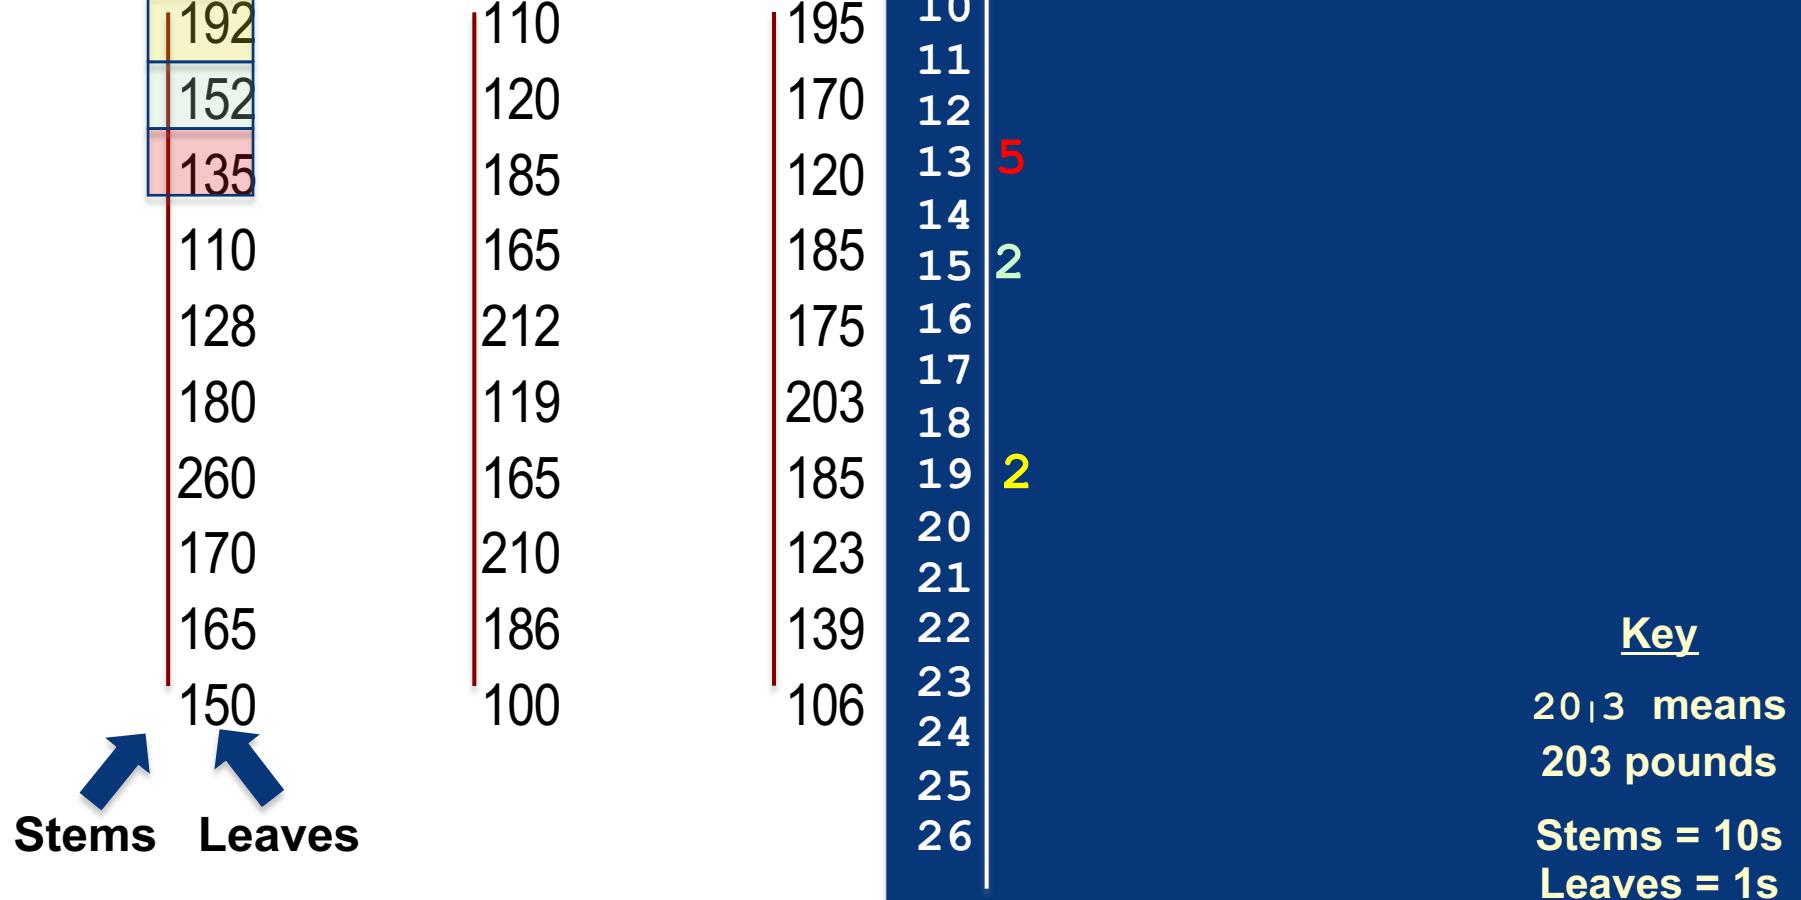
To construct a stemplot:

- Separate each observation into a **stem** (all but the rightmost digit) and a **leaf** (the remaining digit).
- Write the stems in a vertical column; draw a vertical line to the right of the stems.
- Write each leaf in the row to the right of its stem; order leaves if desired.

Stemplots 2

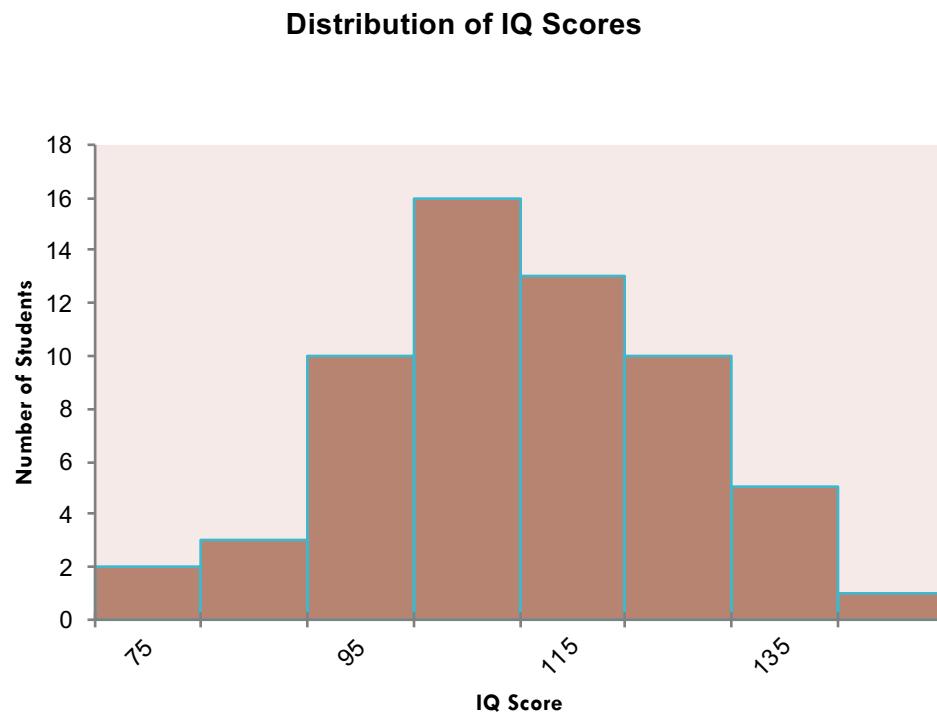
11

Example: Weight Data — Introductory Statistics Class



Histograms

Example: IQ Scores – 60 5th graders



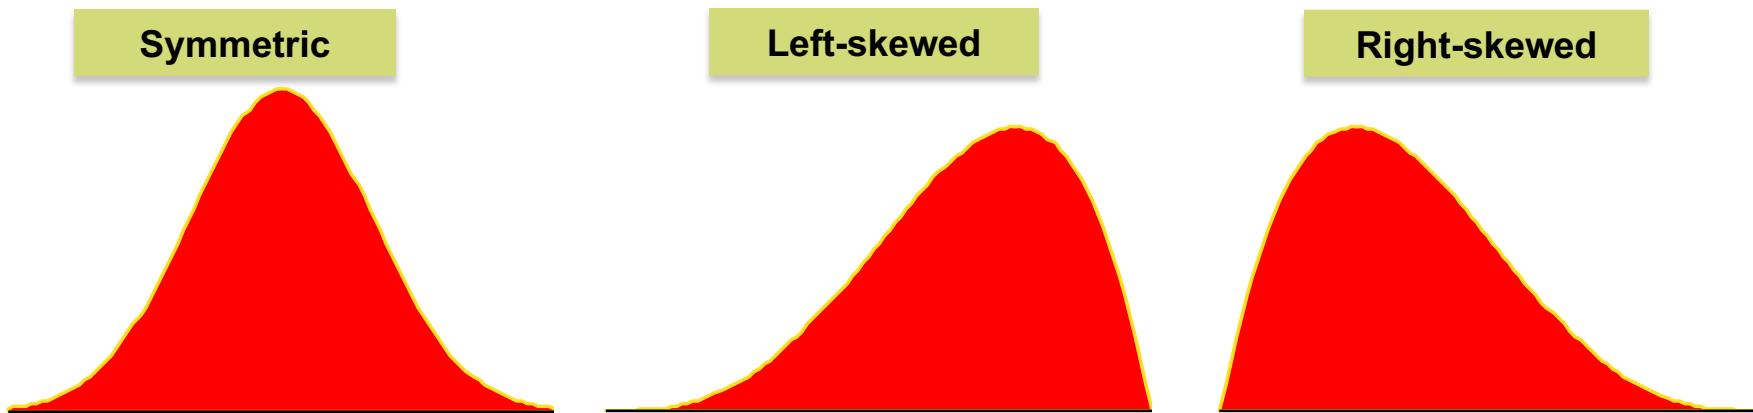
Class	Count
75 ≤ IQ Score < 85	2
85 ≤ IQ Score < 95	3
95 ≤ IQ Score < 105	10
105 ≤ IQ Score < 115	16
115 ≤ IQ Score < 125	13
125 ≤ IQ Score < 135	10
135 ≤ IQ Score < 145	5
145 ≤ IQ Score < 155	1

Examining Distributions 1

- In any graph of data, look for the **overall pattern** and for striking **deviations** from that pattern.
- You can describe the overall pattern by its **shape**, **center**, and **spread**.
- An important kind of deviation is an **outlier**, an individual that falls outside the overall pattern.

Examining Distributions 2

- A distribution is **symmetric** if the right and left sides of the graph are approximately mirror images of each other.
- A distribution is **skewed to the right** (right-skewed) if the right side of the graph (containing the half of the observations with larger values) is much longer than the left side.
- It is **skewed to the left** (left-skewed) if the left side of the graph is much longer than the right side.

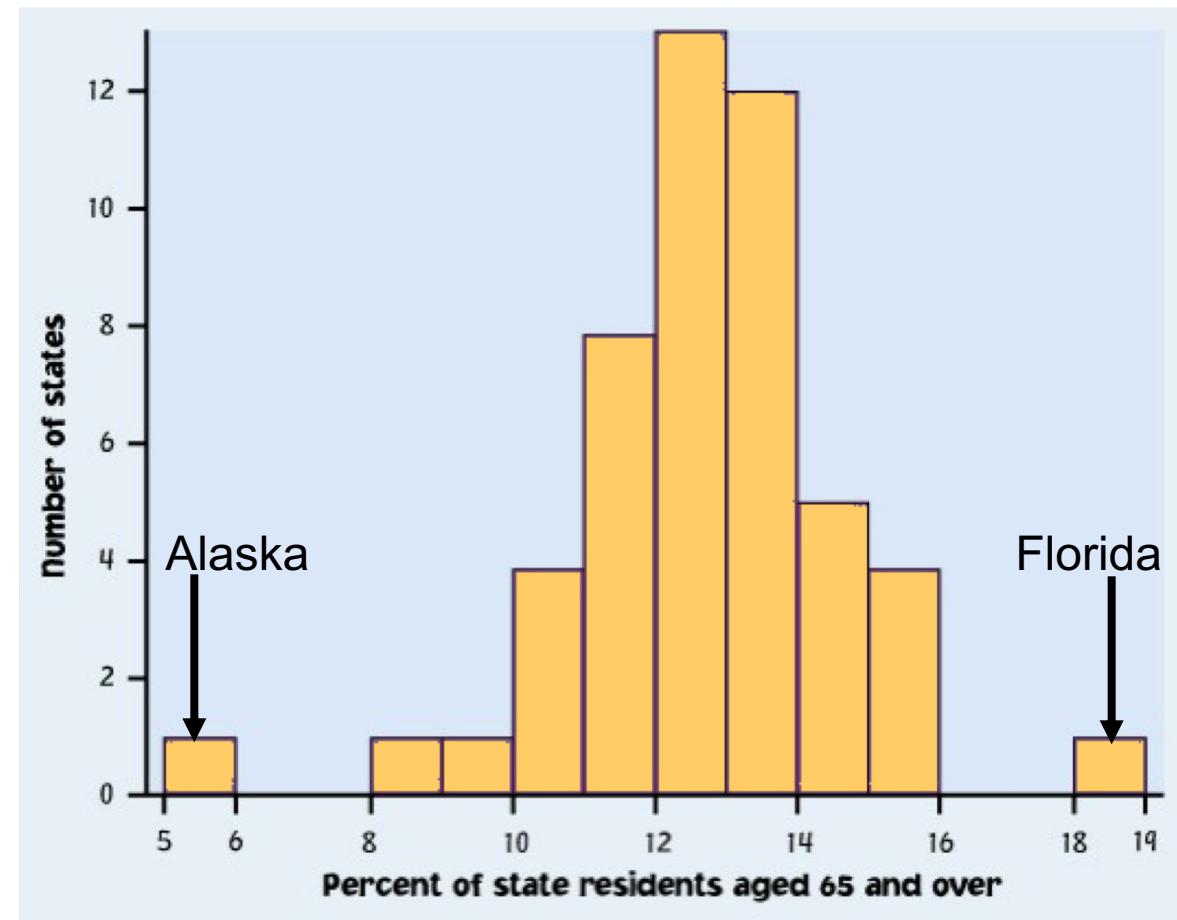


Outliers

An important kind of deviation is an **outlier**. Outliers are observations that lie outside the overall pattern of a distribution. Always look for outliers and try to explain them.

The overall pattern is fairly symmetrical except for two states that clearly do not belong to the main pattern. Alaska and Florida have unusually small and large percents, respectively, of elderly residents in their populations.

A large gap in the distribution is typically a sign of an outlier.



1.3 Describing Distributions with Numbers

- Measuring center: mean
- Measuring center: median
- Measuring spread: quartiles
- Five-number summary and boxplot
- Measuring spread: standard deviation
- Choosing among summary statistics
- Changing the unit of measurement

Measuring Center: The Mean

The most common measure of center is the arithmetic average, or **mean**.

To find the **mean** \bar{x} (pronounced “x-bar”) of a set of observations, add their values and divide by the number of observations. If the n observations are $x_1, x_2, x_3, \dots, x_n$, their mean is

$$\bar{x} = \frac{\text{sum of observations}}{n} = \frac{x_1 + x_2 + \cdots + x_n}{n}$$

In a more compact notation:

$$\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i$$

Measuring Center: The Median

Because the mean cannot resist the influence of extreme observations, it is not a **resistant** or **robust** measure of center.

Another common measure of center is the **median**.

The **median M** is the midpoint of a distribution, the number such that half of the observations are smaller and the other half are larger.

To find the median of a distribution:

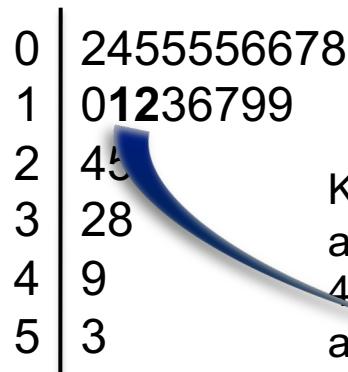
1. Arrange all observations from smallest to largest.
2. If the number of observations n is odd, the median M is the center observation in the ordered list.
3. If the number of observations n is even, the median M is the average of the two center observations in the ordered list.

Measuring Center: Example

Use the data below to calculate the mean and median of the time to start a business (in days) of 24 randomly selected countries.

16	4	5	6	5	7	12	19	10	2	25	19
38	5	24	8	6	5	53	32	13	49	11	17

$$\bar{x} = \frac{16 + 4 + 5 + 6 + \dots + 11 + 17}{24} = 16.292 \text{ days}$$



Key: 4|9 represents
a country that had a
49-day time to start
a business.

$$M = \frac{11 + 12}{2} = 11.5 \text{ days}$$

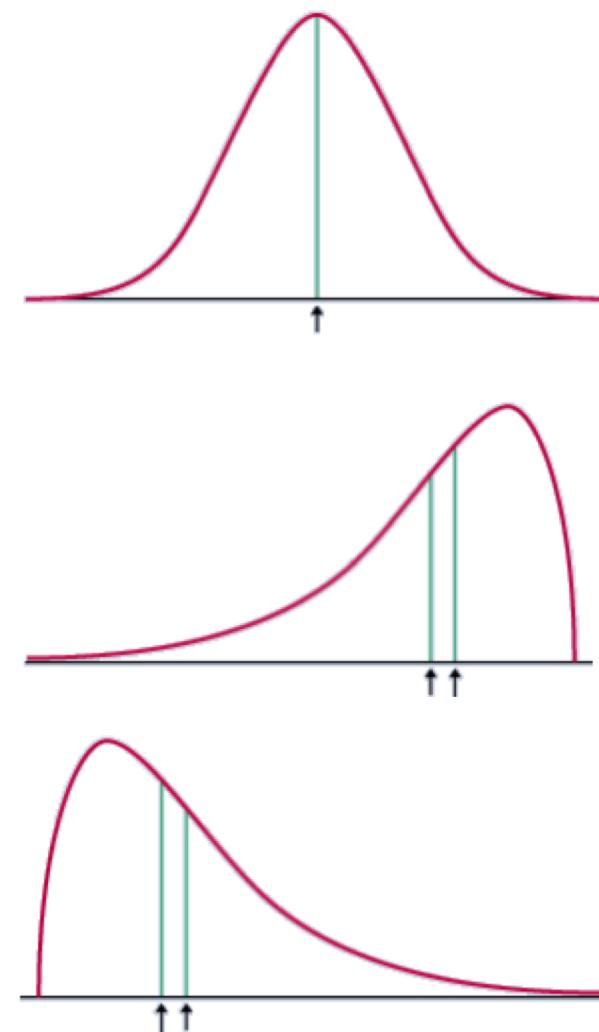
Comparing Mean and Median

The mean and median measure center in different ways, and both are useful.

The mean and median of a roughly **symmetric** distribution are close together.

If the distribution is exactly **symmetric**, the mean and median are exactly the same.

In a **skewed** distribution, the mean is usually farther out in the long tail than is the median.



Measuring Spread: The Quartiles

- A measure of center alone can be misleading.
- A useful numerical description of a distribution requires both a measure of center and a measure of spread.

Calculating the Quartiles

- Arrange the observations in increasing order and locate the **median M** .
- The **first quartile Q_1** is the median of the observations located to the *left* of the median in the ordered list.
- The **third quartile Q_3** is the median of the observations located to the *right* of the median in the ordered list.

The Five-Number Summary

The minimum and maximum values alone tell us little about the distribution as a whole. Likewise, the median and quartiles tell us little about the tails of a distribution.

To get a quick summary of both center and spread, combine all five numbers.

The **five-number summary** of a distribution consists of the smallest observation, the first quartile, the median, the third quartile, and the largest observation, written in order from smallest to largest.

Minimum Q_1 M Q_3 Maximum

Boxplots 1

The median and quartiles divide the distribution roughly into quarters. This leads to a new way to display quantitative data, the **boxplot**.

How to Make a Boxplot

- Draw and label a number line that includes the range of the distribution.
- Draw a central box from Q_1 to Q_3 .
- Note the median M inside the box.
- Extend lines (whiskers) from the box out to the minimum and maximum values that are not outliers.

Suspected Outliers: $1.5 \times \text{IQR}$ Rule

The **interquartile range (IQR)** is defined as $\text{IQR} = Q_3 - Q_1$.

In addition to serving as a measure of spread, the IQR is used as part of a rule of thumb for identifying outliers.

The $1.5 \times \text{IQR}$ Rule for Outliers

Call an observation an outlier if it falls more than $1.5 \times \text{IQR}$ above the third quartile or below the first quartile.

In the business start time data, $Q_1 = 5.5$ days, $Q_3 = 21.5$ days, and so $\text{IQR} = 16$ days.

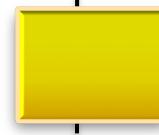
For these data, $1.5 \times \text{IQR} = 1.5(16) = 24$

$$Q_1 - 1.5 \times \text{IQR} = 5.5 - 24 = -18.5$$

$$Q_3 + 1.5 \times \text{IQR} = 21.5 + 24 = 45.5$$

Any business start time shorter than -18.5 days or longer than 45.5 days is considered an outlier.

0	2455556678
1	01236799
2	45
3	28

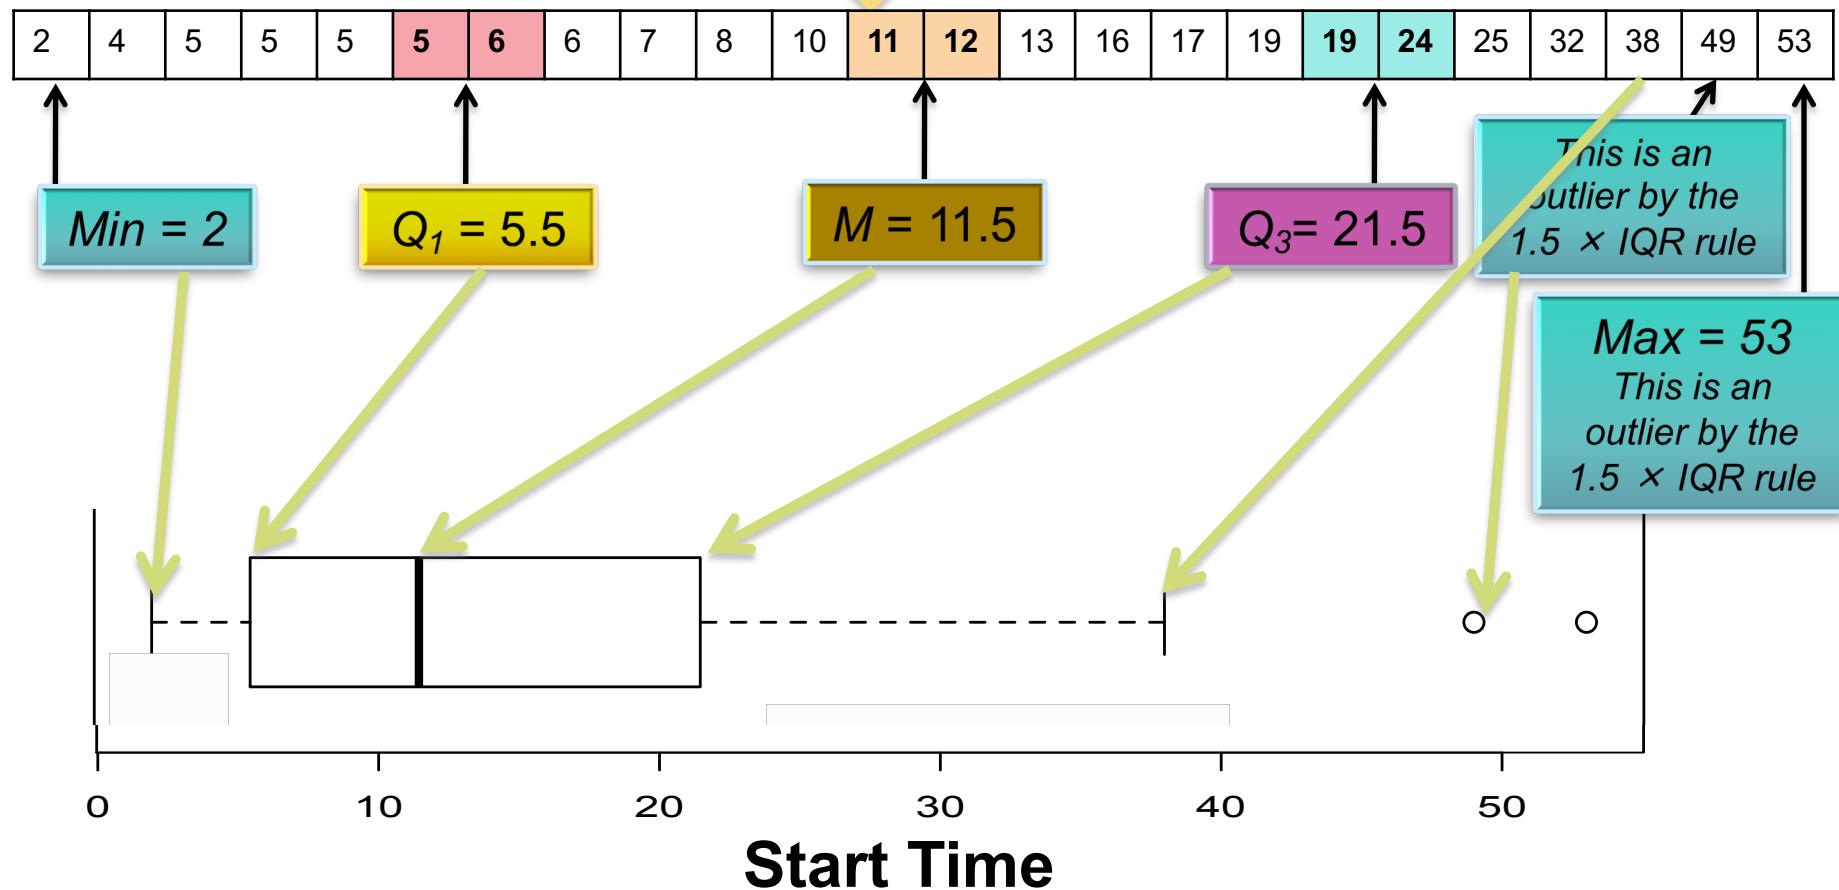


Boxplots 2

Consider our business start times data. Construct a boxplot.

16	4	5	6	5	7	12	19	10	2	25	19
38	5	24	8	6	5	53	32	13	49	11	17

Sort the data



Measuring Spread: The Standard Deviation

The most common measure of spread looks at how far each observation is from the mean. This measure is called the **standard deviation**.

The **standard deviation** s_x measures the average distance of the observations from their mean. It is calculated by finding an average of the squared distances and then taking the square root. This average squared distance is called the **variance**.

$$\text{variance} = s_x^2 = \frac{(x_1 - \bar{x})^2 + (x_2 - \bar{x})^2 + \cdots + (x_n - \bar{x})^2}{n - 1} = \frac{1}{n - 1} \sum (x_i - \bar{x})^2$$

$$\text{standard deviation} = s_x = \sqrt{\frac{1}{n - 1} \sum (x_i - \bar{x})^2}$$

Properties of the Standard Deviation

- s measures spread about the mean and should be used only when the mean is an appropriate measure of center.
- $s = 0$ only when all observations have the same value and there is no spread. Otherwise, $s > 0$.
- s is *not* resistant to outliers.
- s has the same units of measurement as the original observations.

Changing the Unit of Measurement

Variables can be recorded in different units of measurement. Most often, one measurement unit is a **linear transformation** of another measurement unit: $x_{\text{new}} = a + bx$.

Linear transformations do not change the basic shape of a distribution (skew, symmetry, multimodal). But they do change the measures of center and spread.

- Multiplying each observation by a positive number b multiplies both measures of center (mean, median) and spread (IQR, s) by b .
- Adding the same number a (positive or negative) to each observation adds a to measures of center and to quartiles, but it does not change measures of spread (IQR, s).

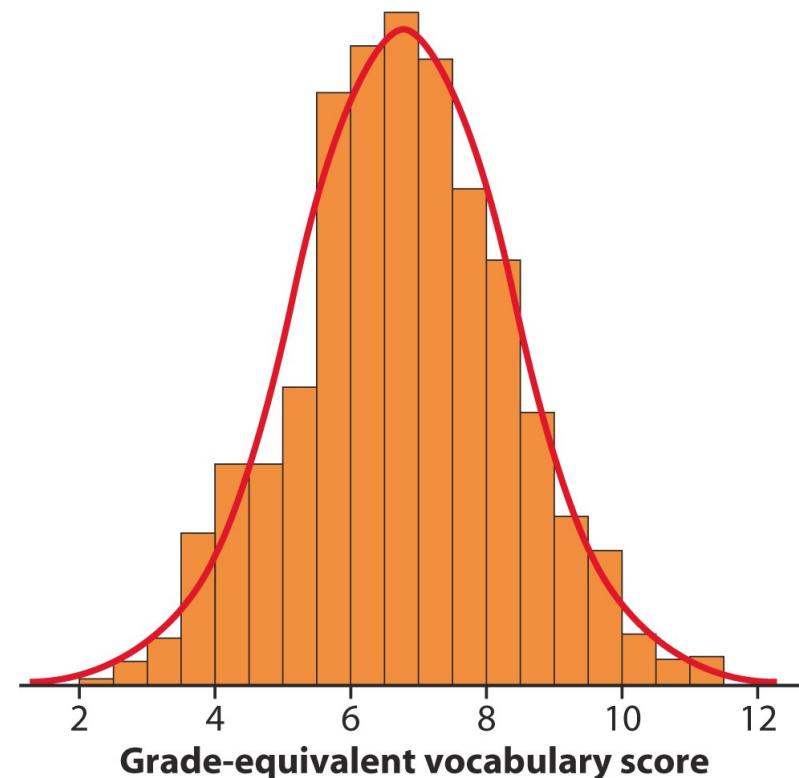
1.4 Density Curves and Normal Distributions

- Density curves
- Measuring center and spread for density curves
- Normal distributions
- Standardizing observations
- Using the standard Normal table
- Inverse Normal calculations
- Normal quantile plots

Density Curves 1

Example: Here is a histogram of vocabulary scores of 947 seventh graders.

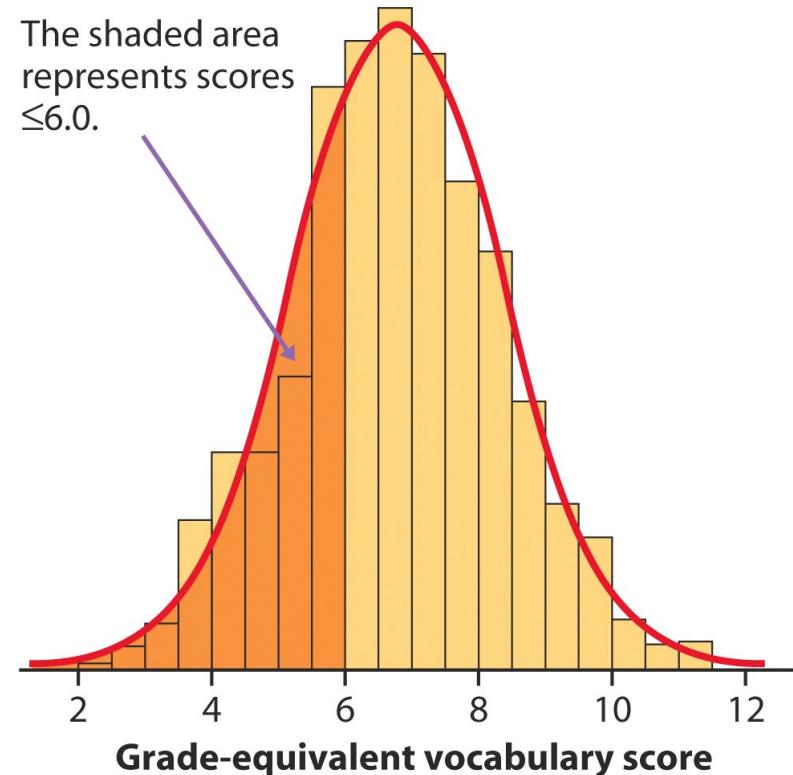
The smooth curve drawn over the histogram is a **mathematical model** for the distribution.



Density Curves 2

The areas of the shaded bars in this histogram represent the proportion of scores in the observed data that are less than or equal to 6.0. This proportion is equal to 0.303.

Now the area under the smooth curve to the left of 6.0 is shaded. If the scale is adjusted so the total area under the curve is exactly 1, then this curve is called a **density curve**. The proportion of the area to the left of 6.0 is now equal to 0.293.



Density Curves 3

A **density curve** is a curve that

- is always on or above the horizontal axis.
- has an area of exactly 1 underneath it.

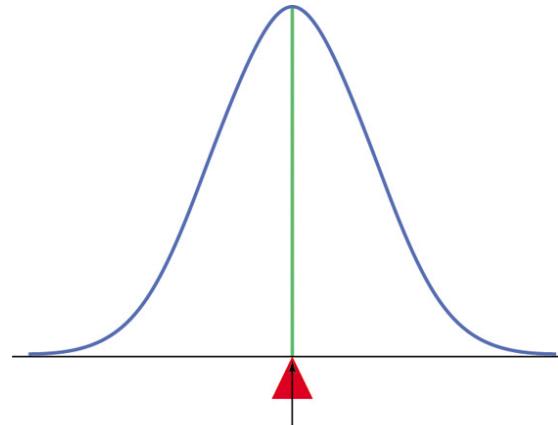
A density curve describes the overall pattern of a distribution. The area under the curve and above any range of values on the horizontal axis is the proportion of all observations that fall in that range.

Density Curves 4

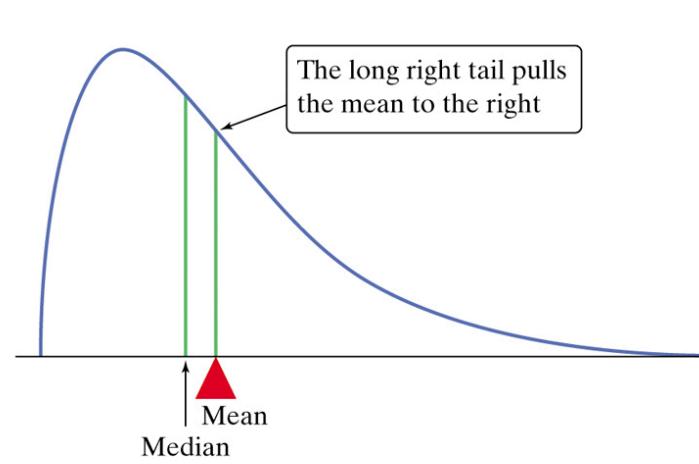
Our measures of center and spread apply to density curves as well as to actual sets of observations.

Distinguishing the Median and Mean of a Density Curve

- The **median** of a density curve is the “equal-areas” point—the point that divides the area under the curve in half.
- The **mean** of a density curve is the balance point, that is, the point at which the curve would balance if made of solid material.
- The median and the mean are the same for a symmetric density curve. They both lie at the center of the curve. The mean of a skewed curve is pulled away from the median in the direction of the long tail.



(a)



(b)

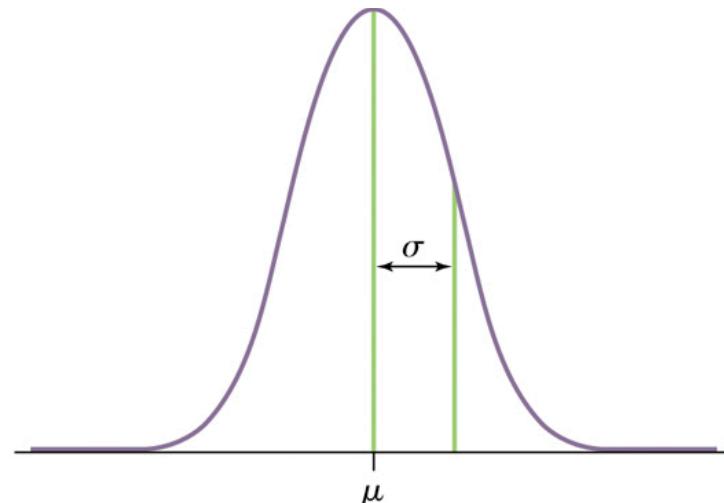
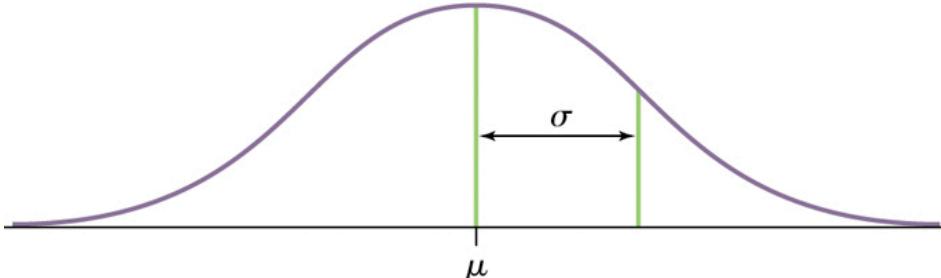
Density Curves 5

- The mean and standard deviation computed from actual observations (data) are denoted by \bar{x} and s , respectively.
- The mean and standard deviation of the actual distribution represented by the density curve are denoted by μ (“mu”) and σ (“sigma”), respectively.

Normal Distributions 1

One particularly important class of density curves is the class of Normal curves, which describe Normal distributions.

- All Normal curves are symmetric, single-peaked, and bell-shaped.
- A specific Normal curve is described by giving its mean μ and standard deviation σ .



Normal Distributions 2

A **Normal distribution** is described by a Normal density curve. Any particular Normal distribution is completely specified by two numbers: its mean μ and standard deviation σ .

- The mean of a Normal distribution is the center of the symmetric **Normal curve**.
- The standard deviation is the distance from the center to the change-of-curvature points on either side.
- We abbreviate the Normal distribution with mean μ and standard deviation σ as $N(\mu, \sigma)$.

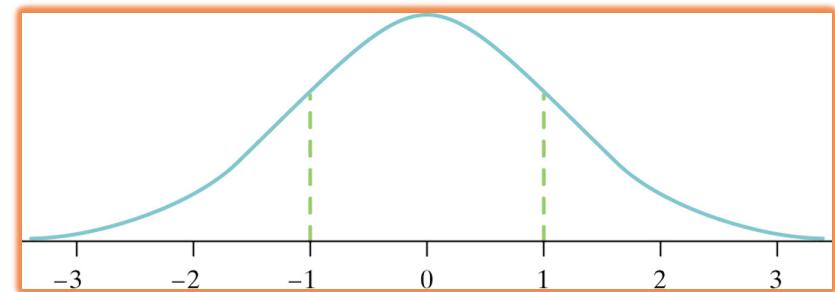
Standardizing Observations

If a variable x has a distribution with mean μ and standard deviation σ , then the **standardized value** of x , or its **z-score**, is

$$z = \frac{x - \mu}{\sigma}$$

All Normal distributions are the same if we measure in units of size σ from the mean μ as center.

The **standard Normal distribution** is the Normal distribution with mean 0 and standard deviation 1. That is, the standard Normal distribution is $N(0,1)$.



The Standard Normal Table 1

Because all Normal distributions are the same when we standardize, we can find areas under any Normal curve from a single table.

The Standard Normal Table

Table A is a table of areas under the standard Normal curve. The table entry for each value z is the area under the curve to the left of z .

TABLE A Standard normal probabilities

z	.00	.01	.02	.03	.04	.05	.06	.07	.08	.09
-3.4	.0003	.0003	.0003	.0003	.0003	.0003	.0003	.0003	.0003	.0002
-3.3	.0005	.0005	.0005	.0004	.0004	.0004	.0004	.0004	.0004	.0003
-3.2	.0007	.0007	.0006	.0006	.0006	.0006	.0006	.0005	.0005	.0005
-3.1	.0010	.0009	.0009	.0009	.0008	.0008	.0008	.0008	.0007	.0007
-3.0	.0013	.0013	.0013	.0012	.0012	.0011	.0011	.0011	.0010	.0010
-2.9	.0019	.0018	.0018	.0017	.0016	.0016	.0015	.0015	.0014	.0014
-2.8	.0026	.0025	.0024	.0023	.0023	.0022	.0021	.0021	.0020	.0019
-2.7	.0035	.0034	.0033	.0032	.0031	.0030	.0029	.0028	.0027	.0026
-2.6	.0047	.0045	.0044	.0043	.0041	.0040	.0039	.0038	.0037	.0036
-2.5	.0062	.0060	.0059	.0057	.0055	.0054	.0052	.0051	.0049	.0048
-2.4	.0082	.0080	.0078	.0075	.0073	.0071	.0069	.0068	.0066	.0064
-2.3	.0107	.0104	.0102	.0099	.0096	.0094	.0091	.0089	.0087	.0084
-2.2	.0139	.0136	.0132	.0129	.0125	.0122	.0119	.0116	.0113	.0110
-2.1	.0179	.0174	.0170	.0166	.0162	.0158	.0154	.0150	.0146	.0143
-2.0	.0228	.0222	.0217	.0212	.0207	.0202	.0197	.0192	.0188	.0183
-1.9	.0287	.0281	.0274	.0268	.0262	.0256	.0250	.0244	.0239	.0233
-1.8	.0359	.0351	.0344	.0336	.0329	.0322	.0314	.0307	.0301	.0294
-1.7	.0446	.0436	.0427	.0418	.0409	.0401	.0392	.0384	.0375	.0367
-1.6	.0548	.0537	.0526	.0516	.0505	.0495	.0485	.0475	.0465	.0455
-1.5	.0668	.0655	.0643	.0630	.0618	.0606	.0594	.0582	.0571	.0559
-1.4	.0808	.0793	.0778	.0764	.0749	.0735	.0721	.0708	.0694	.0681

The Standard Normal Table 2

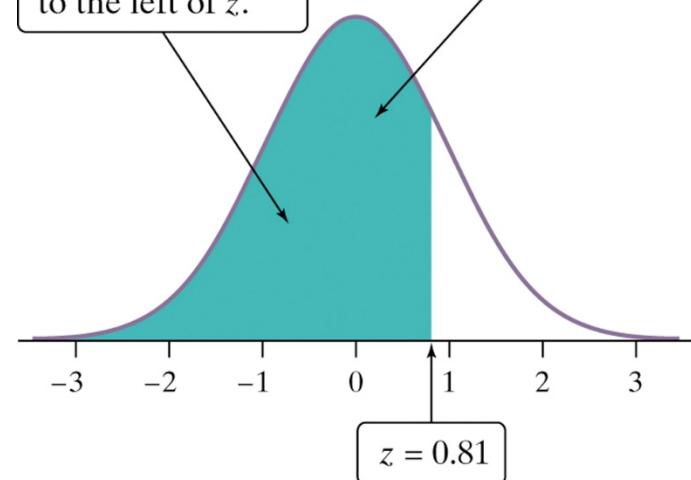
Suppose we want to find the proportion of observations from the standard Normal distribution that are less than 0.81.
We can use Table A.

Z	.00	.01	.02
0.7	.7580	.7611	.7642
0.8	.7881	.7910	.7939
0.9	.8159	.8186	.8212

$$P(z < 0.81) = .7910$$

Table entry for z is always the area under the curve to the left of z .

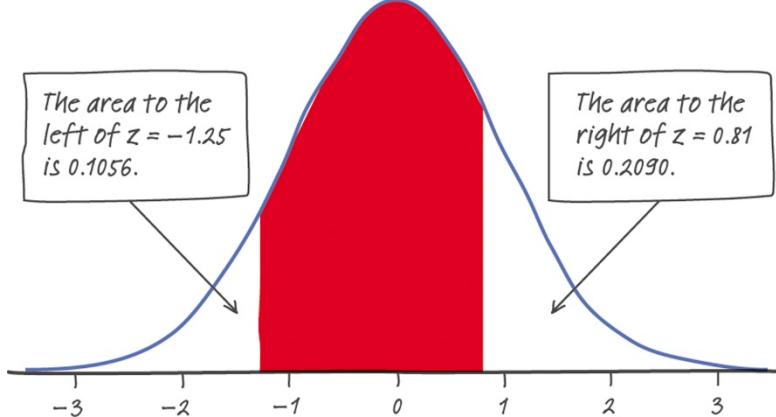
Table entry = 0.7910 for $z = 0.81$.



Normal Calculations 1

Find the proportion of observations from the standard Normal distribution that are between -1.25 and 0.81 .

Can you find the same proportion using a different approach?



$$\begin{aligned}1 - (0.1056 + 0.2090) &= 1 - 0.3146 \\&= \mathbf{0.6854}\end{aligned}$$

Normal Calculations 2

How to Solve Problems Involving Normal Distributions

Express the problem in terms of the observed variable x .

Draw a picture of the distribution and shade the area of interest under the curve.

Perform calculations.

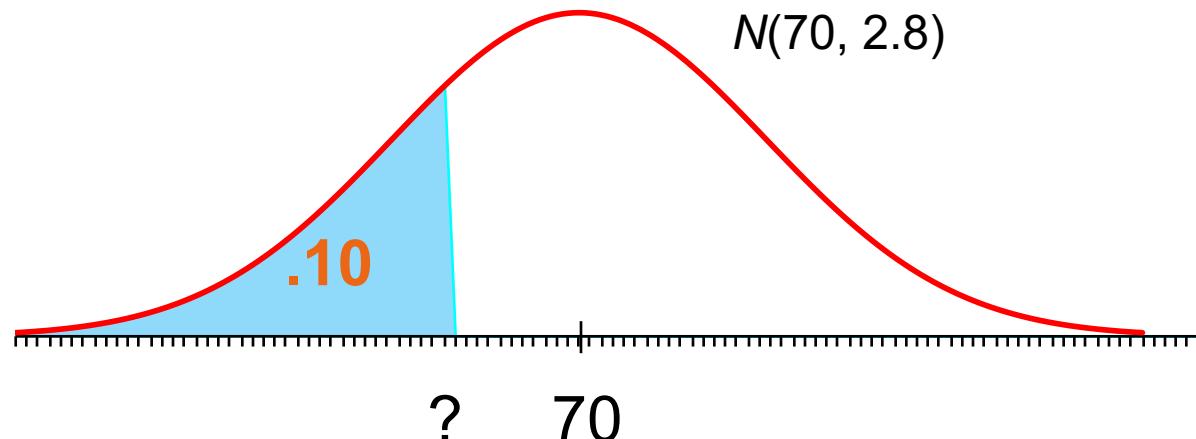
- **Standardize** x to restate the problem in terms of a standard Normal variable z .
- **Use Table A** and the fact that the total area under the curve is 1 to find the required area under the standard Normal curve.

Write your conclusion in the context of the problem.

Normal Calculations 3

According to the Health and Nutrition Examination Study of 1976–1980, the heights (in inches) of adult men aged 18–24 are $N(70, 2.8)$.

If exactly 10% of men aged 18–24 are shorter than a particular man, how tall is he?



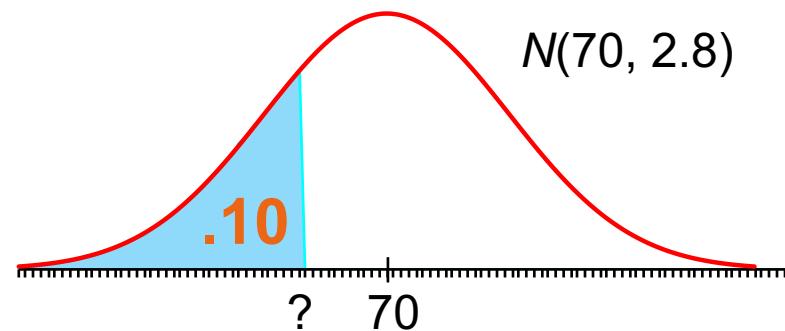
Normal Calculations 4

How tall is a man who is taller than exactly 10% of men aged 18–24?

Look up the probability closest to 0.10 in the table.

Find the corresponding **standardized score**.

The value you seek is that many standard deviations from the mean.



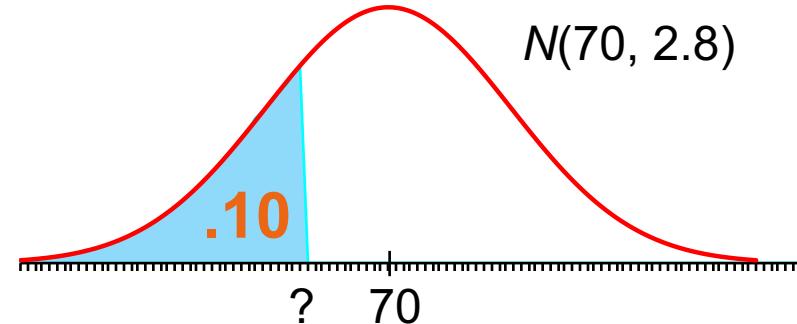
z	.07	.08	.09
-1.3	.0853	.088	.0823
-1.2		.1003	.0985
-1.1	.1210	.1190	.1170

$$Z = -1.28$$

Normal Calculations 5

How tall is a man who is taller than exactly 10% of men aged 18–24?

$$Z = -1.28$$



We need to “unstandardize” the z-score to find the observed value (x):

$$z = \frac{x - \mu}{\sigma} \Rightarrow x = \mu + z\sigma$$

$$\begin{aligned} x &= 70 + z(2.8) \\ &= 70 + [(-1.28) \times (2.8)] \\ &= 70 + (-3.58) = \underline{\underline{66.42}} \end{aligned}$$

A man would have to be approximately 66.42 inches tall or less to be in the lower 10% of all men in the population.

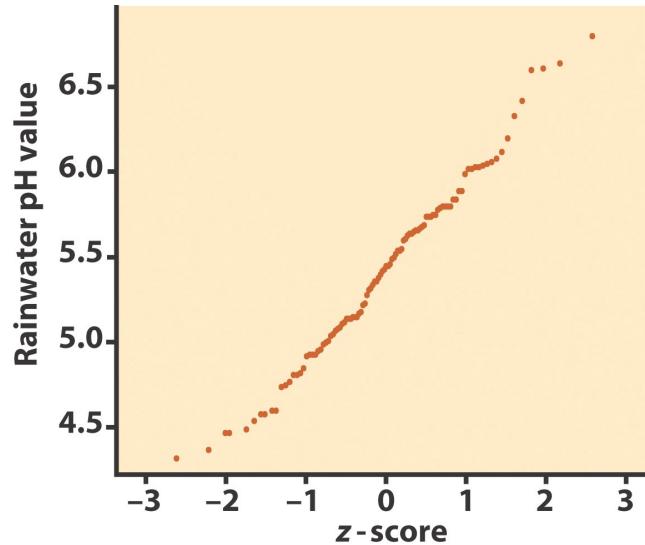
Normal Quantile Plots 1

One way to assess if a distribution is indeed approximately Normal is to plot the data on a **Normal quantile plot**.

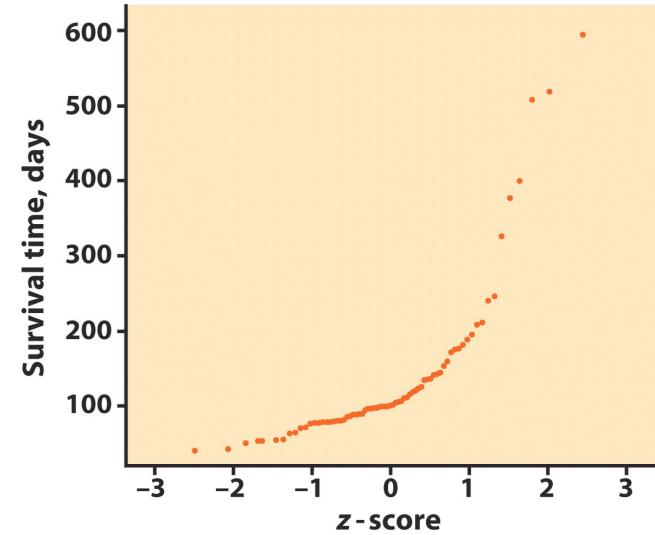
The data points are ranked and the percentile ranks are converted to z-scores with Table A. The z-scores are then used for the x axis against which the data are plotted on the y axis of the Normal quantile plot.

- If the distribution is indeed Normal, the plot will show a straight line, indicating a good match between the data and a Normal distribution.
- Systematic deviations from a straight line indicate a non-Normal distribution. Outliers appear as points that are far away from the overall pattern of the plot.

Normal Quantile Plots 2



Good fit to a straight line: The distribution of rainwater pH values is close to Normal.



Curved pattern: The data are not Normally distributed. Instead, the data are right skewed: A few individuals have particularly long survival times.

Normal quantile plots are complex to do by hand, but they are standard features in most statistical software.